

Programming socket-independent network functions with nethuns

Nicola Bonelli, Fabio Del Vigna, Alessandra Fais,
Giuseppe Lettieri, Gregorio Procissi

Agenda

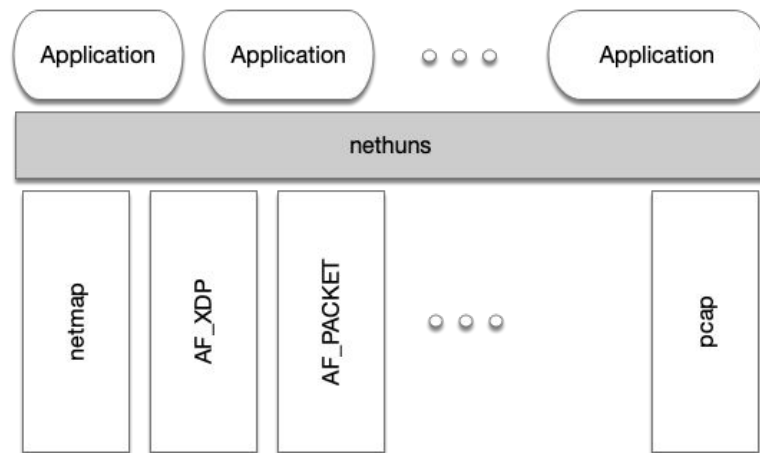
- Motivations
- The `nethuns` architecture
- An essential API
- Engine specific discussion
 - `AF_XDP` and `netmap` cases
- Performance
 - Speed tests
- Practical use cases
 - High-speed traffic replay
 - Open vSwitch

Motivations

- Software data planes are very popular and flexible solutions in real deployments
 - need accelerated I/O to scale up with speed
 - netmap, AF_XDP, ..., each of them with its own “grammar” (API)
 - lack of portability
- nethuns: a **unified abstraction layer** against multiple network APIs
 - nethuns is open source and available at <https://github.com/larthia/nethuns>

nethuns at a glance

- nethuns is a user-space library
 - offers a socket-like API implemented by multiple network I/O engines
- Apps built on top of nethuns must select an I/O engine at compile time
 - this allows compile-time specialization/optimization
 - porting the app over another engine/system only requires recompiling the app itself



nethuns general API


- **Socket** abstraction

- points of access to hardware interface queues
- not a “real” socket though

- **Ring** abstraction

- circular queue of packet descriptors
- include packet metadata and an *in-use* flag

- The actual memory for rings and packet buffers is allocated and exposed/disposed *through* the library
- Each socket can have *at most one ring* for direction (TX/RX)



nethuns at work: opening a socket

- The first step is opening a socket
`sock = nethuns_open(opt)`
- The list of options is quite long...
 - General options
 - `rx`, `tx`, both, ring sizes, zero copy, blocking/non-blocking, etc...
 - Engine specific options
 - name of the eBPF program to inject in the kernel for preprocessing

nethuns at work: binding the socket

- The socket must be bound to a network device
 - A NIC or a specific hw queue of a NIC
`nethuns_bind(sock, dev, queue)`
- Finally, to close the socket
`nethuns_close(opt)`

Receiving packets

```
pktid = nethuns_recv(sock, &pkthdr, &pkt)
```

- Returns
 - the identifier to the next packet in the stream
 - two pointers to
 - a packet header containing metadata
 - a **buffer** containing the full packet
- Once data has been consumed, the buffer must be returned to the library through
`nethuns_release(sock, pktid)`


Nethuns also provides an optimized cache-friendly SPSC queue for dispatching packets to multiple workers!

Transmitting packets

- Packet is queued for transmission by
`nethuns_send(sock, pkt, size)`
- To wake up transmission, need to call
`nethuns_flush(sock)`

nethuns and AF_XDP (1)

- When a nethuns socket is opened, the library creates
 - rx and/or tx rings
 - the umem, with buffers permanently mapped to the corresponding nethuns slot



- **nethuns_recv()**
 - scan for released slots and push them in the fill queue
 - scan for freshly arrived frames in the rx queue and mark the corresponding ring slot accordingly
- Similar behavior is performed on transmission by the **nethuns_flush()**

nethuns and AF_XDP (2)

- In the RX path, AF_XDP uses the eBPF/XDP infrastructure to divert packets from their “canonical” path (through the kernel)
 - need to download and attach to the XDP hook an eBPF program to redirect each packet
 - nethuns default XDP program diverts all packets to AF_XDP (at user space) and manages the kernel maps for opening/closing new sockets
- User defined eBPF program can be loaded by specifying its name in the options upon socket opening
- The `nethuns_close()` function unloads the XDP program attached to the interface

nethuns and netmap

- `netmap` abstracts hw rings through “`netmap` rings”
 - the number of slots in a `netmap` ring is bound to the HW queue size
 - each ring contains pointers to `netmap` buffers (which contain the full packet) shared with the hardware
- The number of `netmap` buffers is configurable (so-called “extra buffers”)
 - use this feature to allocate a `netmap` buffer for each `nethuns` ring
 - allow to decouple `nethuns` rings from `netmap` rings

Performance Evaluation

- `nethuns` is an overlay to standard and accelerated network I/O sockets
 - need to access its performance overhead (if any)
 - very simple testbed: sender/receiver 8 core XEON servers running native and *nethunized* apps... let's see what happens
 - `nethuns-send` and `nethuns-meter` for TX/RX
 - `pkt-gen` (`netmap`) and `xdpsock` (`AF_XDP`)



Speed test: packet transmission

pkt-gen (netmap)
xdpsock (AF_XDP)
nethuns-send




10/40 Gbps

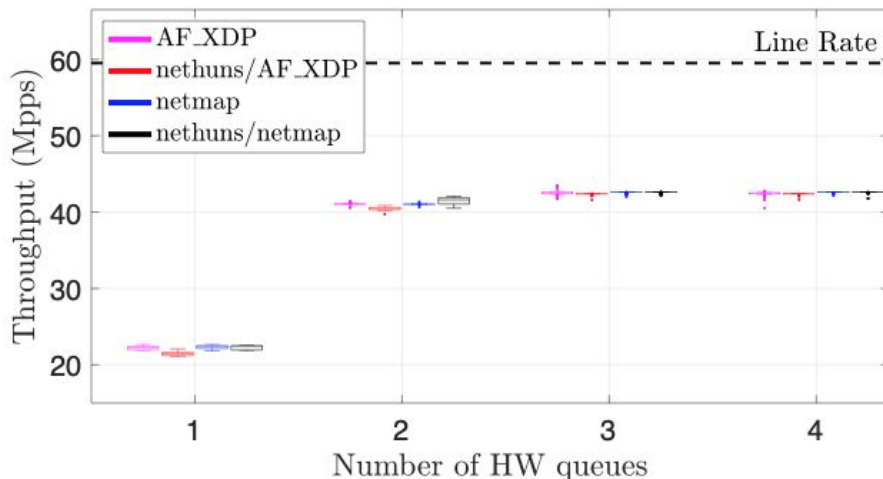


pkt-gen (netmap)

10 Gbps



40 Gbps



Speed test: packet capture

pkt-gen (netmap)

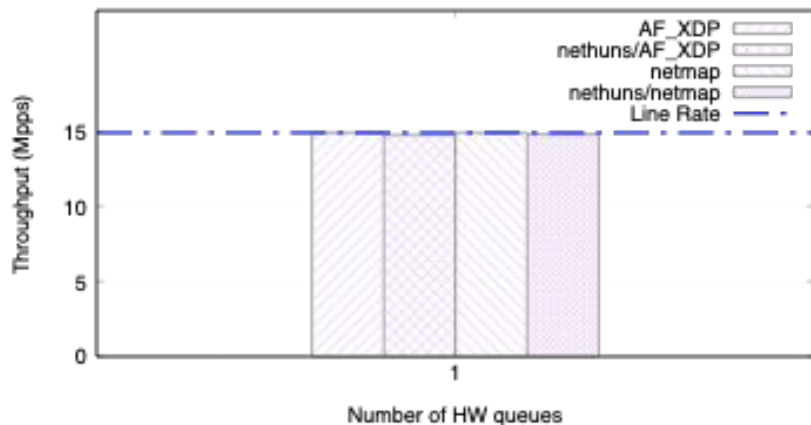


10/40 Gbps

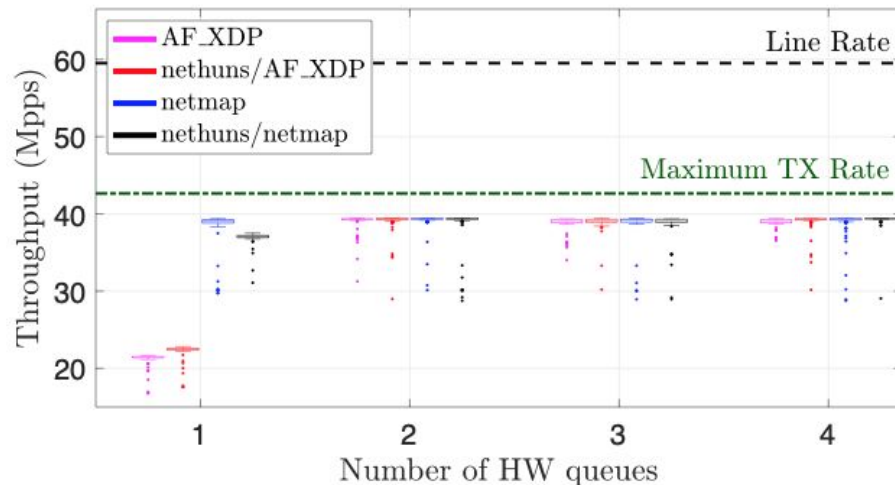


pkt-gen (netmap)
xdpsock (AF_XDP)
nethuns-meter

10 Gbps



40 Gbps



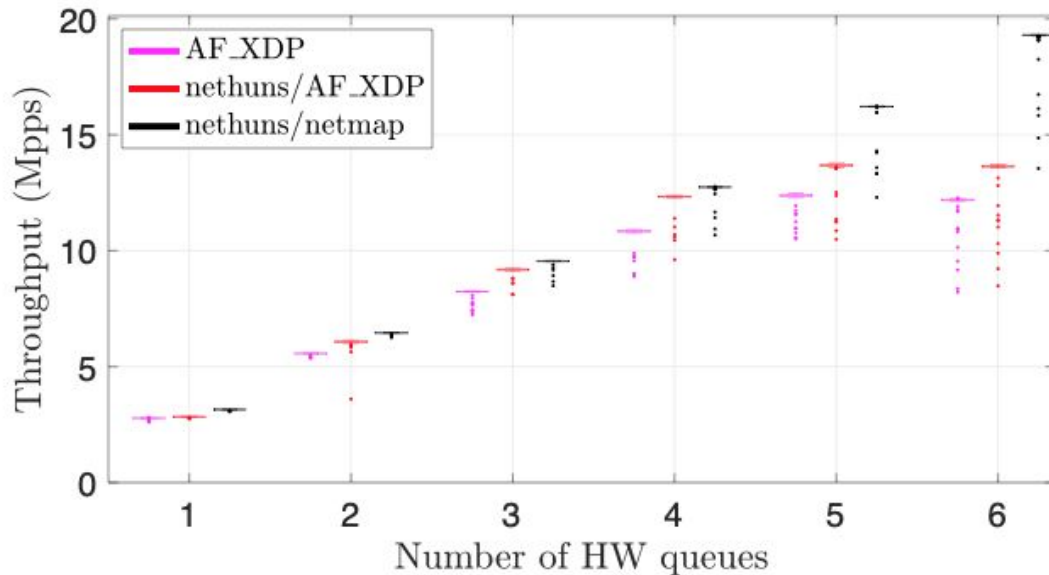

Use case: traffic generator

- `nmreplay` (from the `netmap` ecosystem) replays packets from a `pcap` trace at configurable speed
 - reaches around 21 Gbps speed over a 40Gbit link when playing a specific `pcap` dump
- ported the code on top of `nethuns`
 - replace the `netmap` API calls to those of `nethuns`
 - quick and painless... in around an hour of work we got a traffic generator for `AF_XDP` too!
- **same performance** with `nethuns` over both `netmap` and `AF_XDP`

Use case: Open vSwitch (1)

- user space datapath of OVS uses a generic netdev interface
 - existing netdevs: DPDK, AF_XDP
- Implemented a new `nethuns` netdev
 - the AF_XDP is much simpler (complexity buried within `nethuns`)
 - around 700 lines of code change
 - got OVS for `netmap` “for free”
- And performance?

Use case: Open vSwitch (2)



Conclusion

- **nethuns** is user space library that provides a **socket independent** abstraction to network programming
- currently supported:
 - `AF_PACKET`, `libpcap`, `netmap`, `AF_XDP`
- **porting** an application over a different I/O framework requires **a simple recompilation**
- **no performance degradation** w.r.t. the use of native sockets
- `nethuns` is completely **open source**! Help yourself at:
<https://github.com/larthia/nethuns>