




Towards Scalable and Expressive Stream Packet Processing



Relatore: Alessandra Fais



Introduction and Motivations

Introduction

- Modern networks are *softwarized*
 - Specialized hardware replaced by general purpose devices
 - Network functions implemented as software applications
- And they are shared
 - Accommodate a variety of services on the same infrastructure
 - Adapt to requests for service (de-)activation
 - Always guarantee Quality of Service (QoS) requirements
- Network operators need tools that ease the task of network programming
 - Rapid and easy network (re-)configuration and management
 - Continuous real-time monitoring for detecting security issues and performance degradation



A new era of programmable networks

- In-network data plane processing
 - Programmable software switches
 - Programmable abstractions (e.g., P4)
- Commodity hardware for data plane operations (general-purpose servers)
 - Cheaper alternatives to specialized hardware
 - High flexibility
 - Good maturity level
 - Multi-core architectures to enable parallel computations
 - Multi-queue NICs for parallel tapping of packets from the wire
 - Software accelerated packet capture frameworks



The need to process live packet streams on network end hosts

- Perform accurate packet analytic tasks in end-host servers depends on
 - Ability to capture all the packets arriving to the NIC at speeds of 10+ Gbps
 - Design processing applications to handle incoming data fast enough not to lose packets

STRENGTHS:

- Great flexibility, fully programmable platforms
- Multi-gigabit packet handling with huge volumes of traffic can be achieved
 - Use accelerated packet capture frameworks

PROBLEM:

- How to design network applications to scale with the speed of received data and not get overwhelmed?



The need for parallelism

- Design network applications which take advantage of multiple cores

COMPLEX TASK:

- Network programmers must build their applications from scratch by
 - Implementing the interfaces towards the lower hardware level
 - Implementing the proper mechanisms for parallel programming



SOLUTION:

- Leverage the power of Data Stream Processing (DaSP) frameworks
 - Manage parallelism in an effective way
 - Offer high-level programming abstractions
- Let network programmers only focus on the application logic



Background: Data Stream Processing (DaSP)

- Useful everywhere there is the need for continuous and real-time analysis of streaming data
 - Networking, smart cities, smart mobility, smart logistics, ...
- DaSP frameworks
 - Simplify the implementation of efficient streaming computations
 - Designed to implement general-purpose applications
 - Offer programming abstractions of multi-stage pipelines
 - Every stage (operator) can be seamlessly parallelized
 - Lower-level details hidden to the programmer



Background: DaSP frameworks

- Mainstream DaSP frameworks
 - Based on the Java Virtual Machine
 - Support implementation of streaming applications on distributed environments
 - Many overheads impact performance
 - Not possible to keep up with common traffic rates in networking scenario
- A recent proposal: WindFlow
 - Targets single multi-core machines with GPU support
 - C++17-based
 - High performance + high expressiveness





Proving the feasibility of using the DaSP approach in the networking scenario


The need for suitable extensions

- The WindFlow DaSP framework
 - Offers a modern design and high performance
 - However, it is conceived to implement general streaming computations
 - It lacks mechanisms for binding source operators to multiple sources of data (NIC hardware queues)
 - Source operators must be tuned for packet capturing and parsing
- We implemented proper optimizations for the networking domain
 - Full support for parallel source entities
 - Each one attached to a different NIC hardware queue
 - Acceleration of source operators
 - Fully generic memory pool of tuples



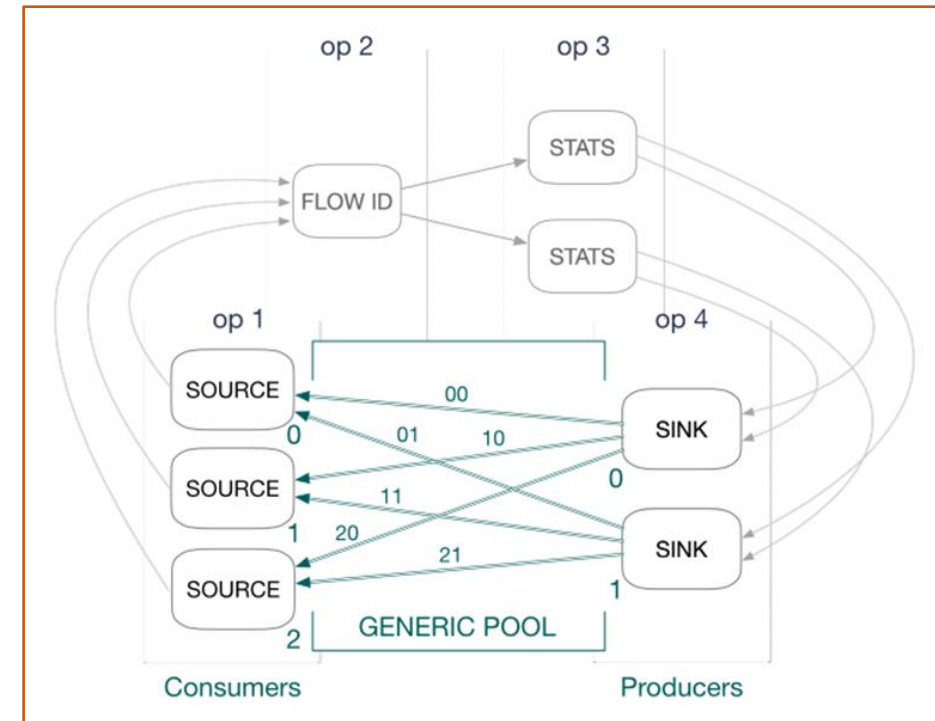
Support for multi-source

- The SOURCE oversees the following tasks
 - Receives captured packets from the NIC
 - Using fast packet capturing tools (e.g., Netmap, AF_XDP, libpcap, AF_PACKET)
 - Module called SNIFFER
 - Parses packet headers and creates accordingly all the tuples that will be processed in the graph
- The SOURCE is the bottleneck in the graph
 - Given its considerable workload
- We removed the performance bottleneck by parallelizing the SOURCE operator




Generic memory pool

- Tuples are the stream elements flowing in the application graph
 - SOURCE nodes initialize tuples from packet contents
 - SINK nodes collect tuples at the end of the processing
- Memory pool and tuple recycling in the application graph to avoid expensive memory allocations
 - Pre-allocated pool of empty tuples
 - Matrix of communication channels
 - Implemented as SPSC* queues
 - Support for generic number of consumer and producers
 - SOURCE consumes tuples from the pool
 - SINK produces tuples into the pool




* SPSC = Single Producer Single Consumer



Experimental evaluation

Raw speed test

- Assess the scalability of the basic framework
- Modules in the system
 - Packet capturing sub-system
 - Inter-thread communication mechanisms
 - Tuple memory pool
- Minimal configuration of the application
 - SOURCE and SINK operators only
 - Tested configurations
 - Chained execution (on the same thread)
 - Non-chained execution (on separated threads)




Raw speed test



Max achievable speed in packet generation around 30 Mpps



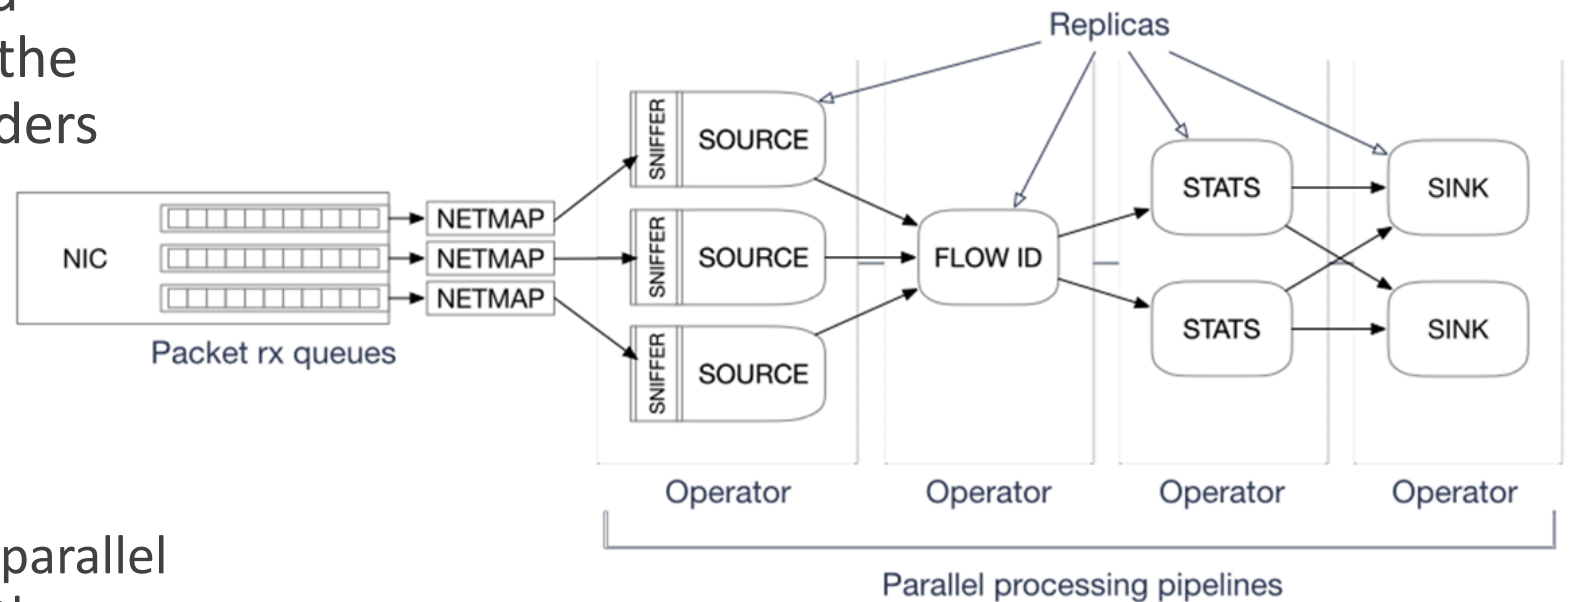
- Workload characteristics
 - Synthetic traffic
 - Minimally sized packets (64 bytes, UDP)
 - Traffic generator runs pkt-gen (Netmap suite)
- Packet rate (millions of packets per second) sustained by the minimal packet-counter application for increasing number of source nodes
- Given X application threads, the number of source nodes in the chaining case is X, while in the no-chaining case is X/2



Flow counter application test

More complex application:

- Produce some packet and flow level analytics from the processing of packet headers



- Pipeline of 4 nodes
 - All can be replicated in parallel independently of the others


Flow counter app test

< / >


Max achievable speed in packet generation around 20 Gbps (~5.2 Mpps)

< / >

- Workload characteristics
 - Realistic traffic
 - pcap file replayed at max speed using nmreplay (Netmap suite)
 - packets of different lengths (generally larger than 64 bytes)
- Packet rate (millions of packets per second) sustained by the test monitoring application with three different combinations of replication and chaining



n-1-1-1	n SOURCE nodes, no other replicated nodes, no chaining
n-1-[1-1]	n SOURCE nodes, no other replicated nodes, STATS node is chained with SINK node
n-1-[2-2]	n SOURCE nodes, one FLOWID node, two nodes each of both STATS and SINK chained pairwise



Conclusions

- App throughput increases linearly with the number of sources
 - Full data elaboration sustained up to nearly 20 Gbps in different configurations
- Proved feasibility of using DaSP approach to compute packet and flow level analytics from live data
 - Common computational pattern for network applications
 - Network monitoring is one of the largest domains of interest
- WindFlow can be a convenient development platform in the networking domain
 - FUTURE DIRECTIONS:
integrate these results into the WindFlow framework to provide a new native networking module with *ad hoc* ready-to-use network-oriented operators