



# Fast and Scalable Network Monitoring Applications based on Data Stream Processing

Presenter: Alessandra Fais

PhD student at University of Pisa



# Outline

- Introduction
  - Problem and main goals
- What can be found in literature
- Proposed approach
  - Continuous monitoring as stream analytics problem
  - Architecture and challenges
- Conclusions



# Introduction

- Modern networks
  - Accommodate a variety of services
  - Shared infrastructure
  - Adapt to requests for service (de-)activation
  - Stringent Quality of Service requirements



# Introduction

- How to manage modern networks?



# Introduction

- How to manage modern networks?
- Network operators need
  - Rapid and easy network (re-)configuration
  - Continuous real-time network monitoring for detecting
    - Security issues
    - Performance degradation

# Introduction

- How to manage modern networks?
- Network operators need
  - Rapid and easy network (re-)configuration
  - Continuous real-time network monitoring for detecting
    - Security issues
    - Performance degradation





# What can be found in literature

- Main tradeoff: **expressiveness** VS **performance**

# What can be found in literature

## ➤ Main tradeoff: **expressiveness** VS **performance**

- Solutions based on

- Stream processing only → high expressiveness
- Programmable switches alone → high performance




# What can be found in literature

## ➤ Main tradeoff: **expressiveness** VS **performance**

- Solutions based on
  - Stream processing only → high expressiveness
  - Programmable switches alone → high performance
- Can we provide both?
  - Some new proposals use stream processing AND programmable switches
  - Proposed approach based on stream processing ONLY

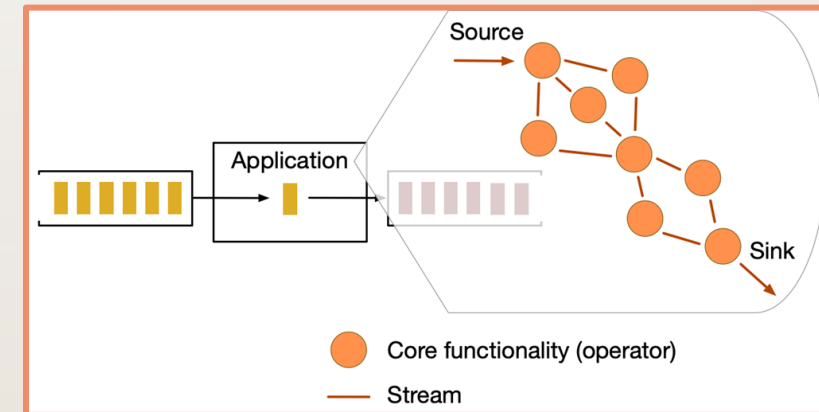
# A new framework: main goals

- **High performance: real-time continuous data analysis**
  - Stream processing pipelines
  - Implementation through Data Stream Processing frameworks




# A new framework: main goals

- **High performance: real-time continuous data analysis**
  - Stream processing pipelines
  - Implementation through Data Stream Processing frameworks
- **Scalability**
  - Distributed processing over programmable network nodes

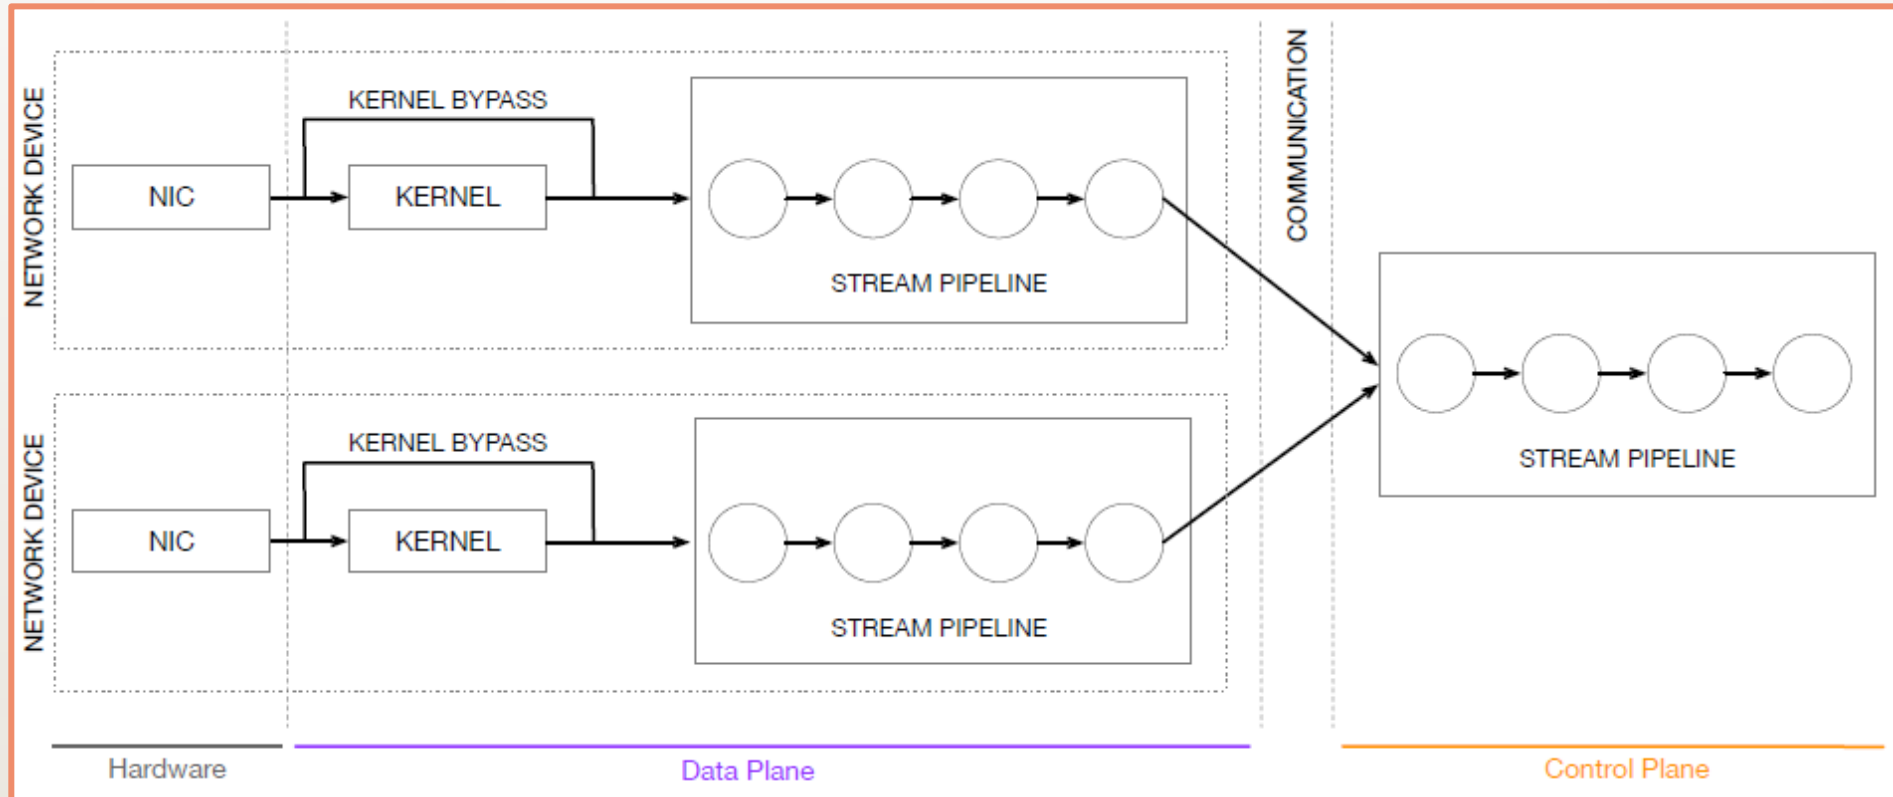


# A new framework: main goals

- **High performance: real-time continuous data analysis**
  - Stream processing pipelines
  - Implementation through Data Stream Processing frameworks
- **Scalability**
  - Distributed processing over programmable network nodes
- **High expressiveness**
  - Provide high level programming abstractions
  - Support for complex analysis queries

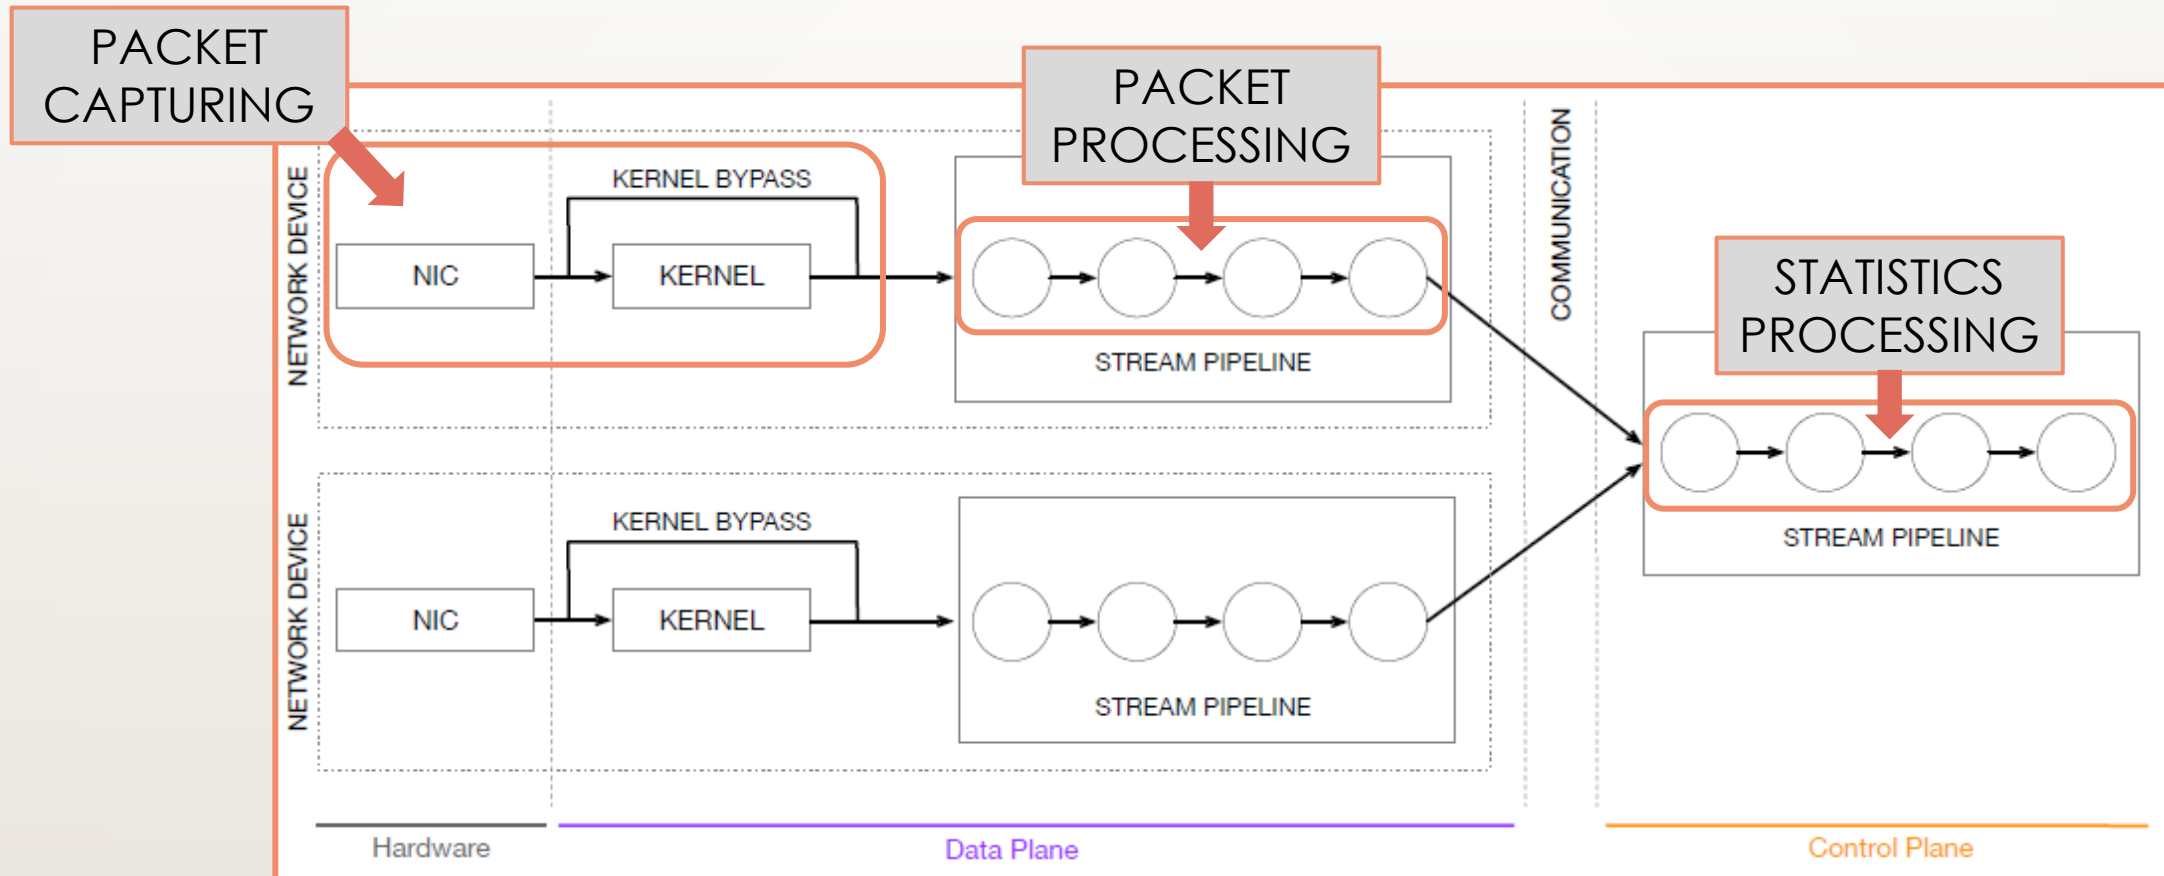


# A new framework: architecture



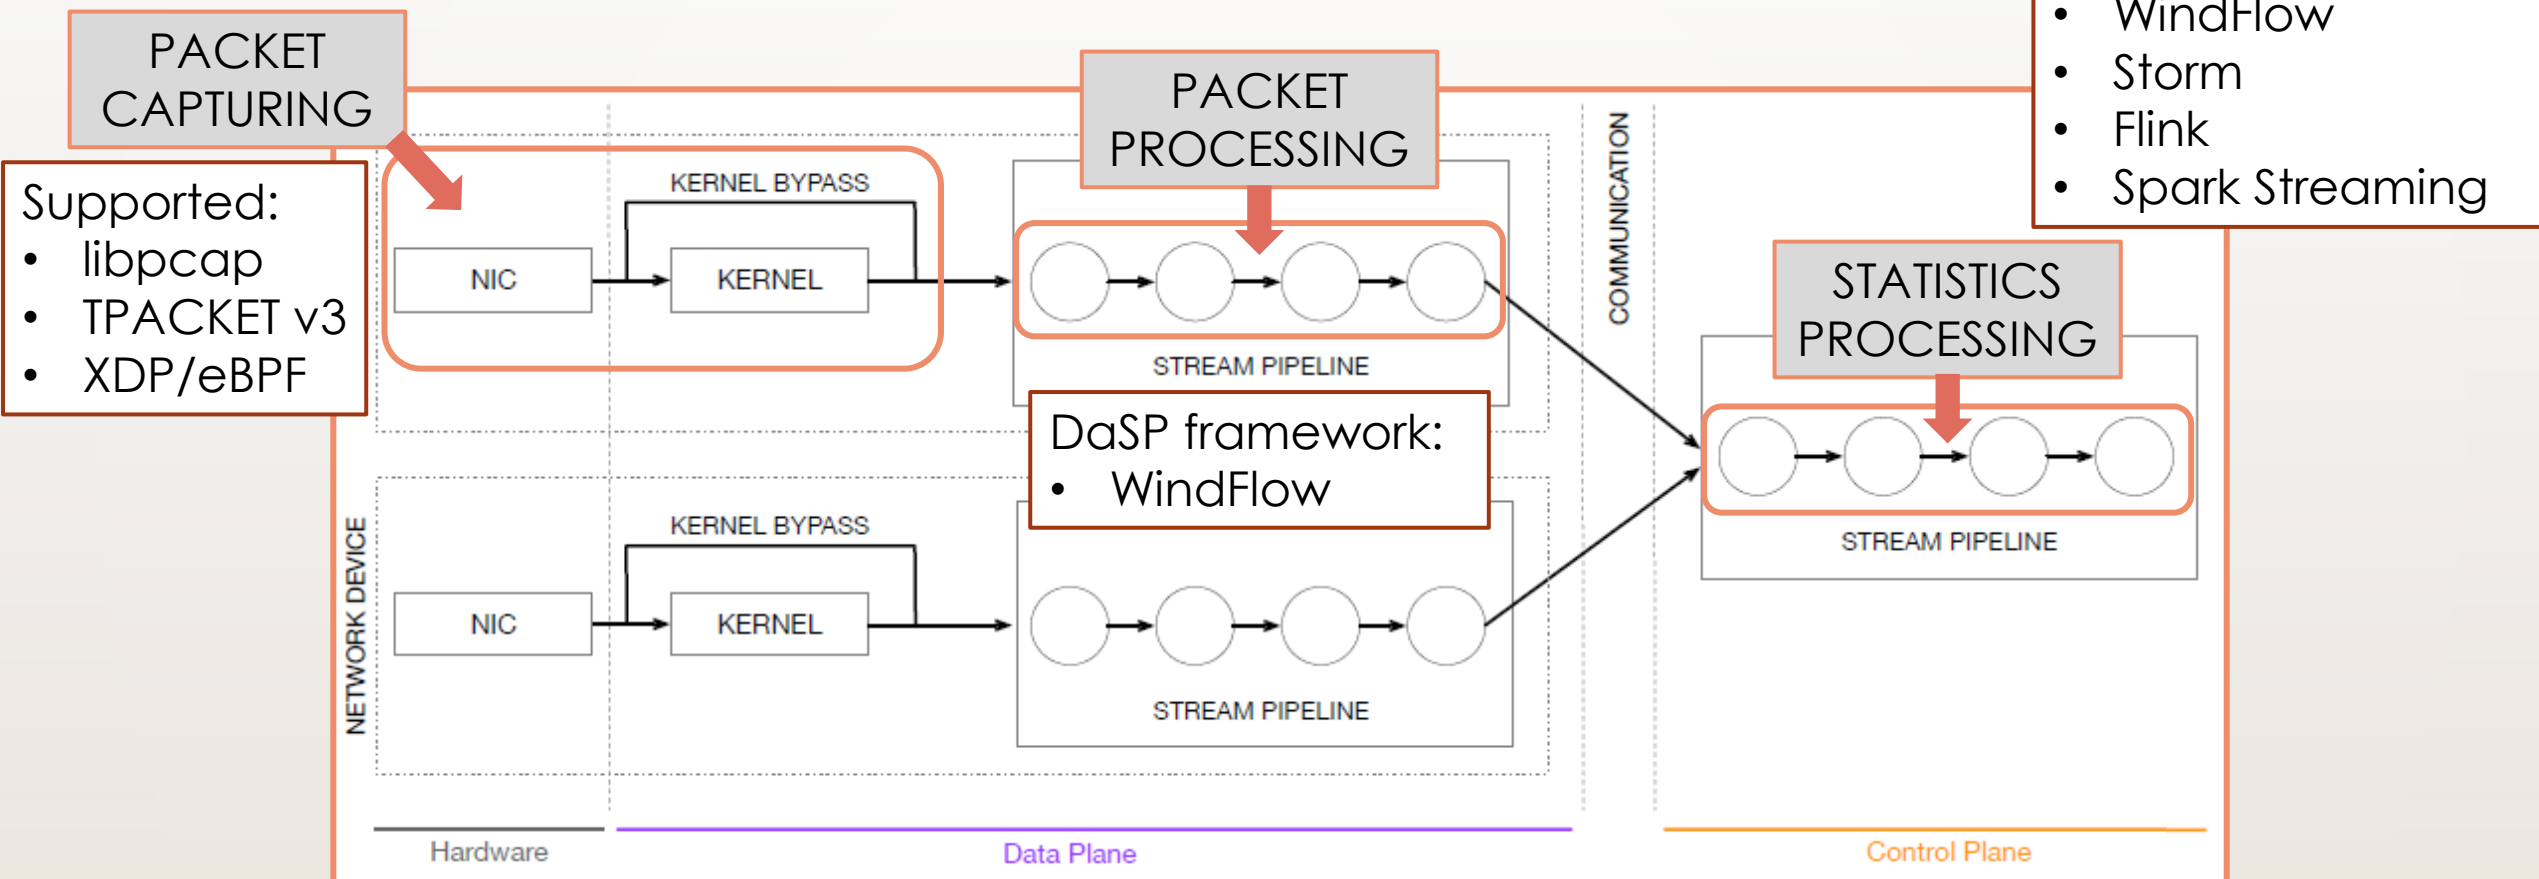


# A new framework: architecture





# A new framework: architecture





# Conclusions

- Cheaper solution
  - Programmable switches are expensive
  - Use general purpose network devices (multicores + SmartNICs, GPUs, FPGAs)

# Conclusions

- Cheaper solution
  - Programmable switches are expensive
  - Use general purpose network devices (multicores + SmartNICs, GPUs, FPGAs)
- High performance
  - Optimized usage of resources
  - Parallelism in all computation stages (packet capturing, processing)



# Conclusions

- Cheaper solution
  - Programmable switches are expensive
  - Use general purpose network devices (multicores + SmartNICs, GPUs, FPGAs)
- High performance
  - Optimized usage of resources
  - Parallelism in all computation stages (packet capturing, processing)
- High-level interface
  - Simplify the implementation of real-time continuous monitoring applications
  - Encourage network programmers to use the proposed framework